


THE FORK IN THE ROAD

LEADERSHIP LESSONS FOR A
WORLD WITH NO PLAYBOOK



MICHAEL DONOSER

ABOUT THESE STORIES

Over eleven years at Amazon—leading cross-functional tech innovation teams building products that reached millions of customers—I made an extraordinary number of mistakes. I missed signals I should have caught. I ran meetings that should never have happened. I gave feedback that landed wrong, made decisions that looked reasonable at the time and embarrassingly stupid in hindsight, and occasionally led a room full of talented people directly into a wall.

What I also did, eventually, was pay attention. I learned from my mistakes, and I doubled down on a natural curiosity—a stubborn need to understand why things work the way they do—that, over time, turned those mistakes into something more useful than regret.

The question that consumed me more than any technical challenge was simpler and stranger: why do some teams seem almost effortlessly effective—energised, honest, driven, resilient under pressure—while others, staffed with equally talented people and similar resources, quietly struggle? I watched this play out repeatedly. Same company. Same tools. Same hiring bar. Completely different outcomes.

My working conclusion, validated over many years and countless situations, is that the difference almost never lives in the skillset. It lives in the culture—in the small, almost invisible choices that accumulate into habits defining how a team actually behaves and communicates, and especially how it holds together when things get hard. As they always do.

I spent a long time thinking about how to share what I had learned. Frameworks felt too clean, principles too abstract. I kept coming back to stories—because that is where these lessons actually live. So instead of laying them out as theory, I captured them as human situations where

the gap between intention and behaviour becomes clear. The stories are hypothetical—not transcripts of real events—but grounded in situations I witnessed, participated in, or caused. I've changed the details, compressed the timelines, and constructed composite characters. The underlying dynamics are real. I've lived most of them.

Each story follows the same structure: a situation any tech company might face, shown through two hypothetical teams—Team Alpha and Team Beta—with no meaningful difference in skill, experience, or intent. The difference is in how they behave, and in the specific mechanisms—repeatable, concrete practices—that produce that behaviour. At the end of each story, I describe those mechanisms directly. Not as values or aspirations, but as things you can start doing on Monday morning.

One more thing worth saying clearly: none of what follows is only for the boss, or for the most senior person in the room. Everyone is a leader. Every person on a team, regardless of title or tenure, can shape culture and influence how the work gets done. The behaviours in these stories are not reserved for those with authority—they are things anyone, at any level, can and should own.

I have no illusion that the ideas in these pages are all original—most of what I found useful I found elsewhere, in the work of people far more rigorous than I am. My contribution is to connect those ideas into something coherent and anchor them in situations where they actually matter. I got a lot wrong over the years. I'm sharing what I learned in the hope that some of it is useful to someone else.

This is the second story.

EXAMPLE CHAPTER

The Olympus Bet

*How two teams handled the mid-year pivot
and the mechanisms that defined their delivery*

THE YEAR'S BIGGEST BET

THE PROJECT had a name that matched its ambitions. “Olympus” was the centrepiece of the year’s strategy—a new product recommendation experience, positioned in the annual planning cycle as a system that would fundamentally transform how customers discovered and engaged with products. The business case was compelling: increased customer engagement, improved product discovery, higher conversion rates, and a measurable uplift in revenue. The promise was not incremental. It was transformational.

After a gruelling planning cycle that stretched through the spring, the goals were locked: Olympus would be delivered by November 30th. The team assigned to build it was well-staffed. The budget was approved. The architecture had been agreed upon in detail across multiple planning sessions. And the expectations from the executive suite—quietly but clearly communicated—were sky-high.

The team went into full execution mode in early April, and the first weeks had the particular energy of a well-resourced project at its beginning: clarity of purpose, full calendars, and the forward momentum that comes from finally building something after months of planning.

From April and into August, the team was heads down. Engineers were deep in their work streams, sprint reviews were busy, and the project boards filled with tickets in progress. Scientists were zeroing in on the algorithm best suited to the product’s requirements. The excitement of Olympus—the scale of what it could be, the visibility it carried—had not worn off. But pressure was beginning to make itself felt. A few threads were more complex than the planning had suggested. A dependency on a third-party data pipeline had slipped two weeks. The velocity felt right, but the margin for error was quietly narrowing.

Everything was broadly on track. Status updates were green. Leadership projected the calm, measured confidence of people who believed the deadline was achievable—tight, perhaps, but achievable. Privately, some engineers had begun doing the mathematics of what remained and wondering whether the buffer built into the plan was really as generous as it had looked in April. But nobody said so, not in any meeting that mattered. There was too much to do, and too much at stake.

By September, however, the world had moved. The shift was technological: new capabilities had emerged in the market—capabilities that did not exist when the team first committed to their architecture. A competitor had moved quickly and demonstrated what could be enabled, and the early results were difficult to ignore. Internal data pointed in the same direction. A different algorithmic and architectural approach—absent from the original design documents—appeared to offer materially better customer outcomes. Not a marginal improvement, but a potentially substantial one.

Nothing about the original plan was wrong. It had been built on the best available information at the time, and the team had executed it well. But new information had arrived, and with it a question that hadn't existed six months earlier: was the path they were on still the right one?

That question arrived at precisely the moment when the team was deepest in the execution of the original plan.

* * *

This was the fork in the road.

From here, the story splits. In one world is Team Alpha. In the other, Team Beta. Both received the same fresh data. Both faced the same situation. What followed was not determined by talent or resources—it

was determined by the culture, habits, and mechanisms each team had built, or failed to build, in the months before that moment arrived.

Let's trace how our two teams Alpha and Beta execute the project in the months that follow.

TEAM ALPHA: THE POTEMKIN GRIND

The Commitment That Would Not Bend

Across Team Alpha, several engineers, scientists, and product staff had independently tracked the same technological shift. They had seen the competitor's results, run their own analysis, and tested the new capabilities against the product's requirements. Their findings were unambiguous: the improvement was real, significant, and directly relevant to what Olympus was trying to achieve. This was not speculation—it was a verified technical conclusion. More than a few had quietly begun to wonder whether the project should pivot.

But they did not say so. Not directly. Not to the leader.

The culture that had formed around him over the preceding months was not one of open challenge. Past experience had taught the team that raising uncomfortable questions had a cost—and the team had learned to edit themselves before they spoke.

So instead of raising the concern individually, a small informal committee formed—people who trusted one another enough to say privately what none would say publicly. They would write a document: thorough, data-driven, airtight.

Drafts circulated. Language was debated. The committee moved slowly, not out of carelessness, but out of fear. And while they polished

and reconsidered, the window for a rapid pivot was quietly narrowing—every week of silence was another week of effort invested in the wrong direction. The cost mounted gradually, almost invisibly, until it became impossible to ignore.

By the time the document had survived three rounds of internal review, the committee's conclusion was clear: Olympus should pivot. The case was solid, the recommendation specific, and the group was certain. What remained was the harder question: who would take it to the leader?

They turned to the most senior product manager on the team. She had the most standing, the longest relationship with Team Alpha's leader, and—in the committee's quiet judgment—the best chance of being heard. She agreed. She asked for time. He agreed, and she sent him the polished document in advance.

The leader of Team Alpha received the document proposing a pivot for Olympus by the end of September. He reviewed it thoroughly and arrived at the meeting with a clear intention: to reject the proposal outright.

This was not indifference—it was the response of someone who had spent eight months pouring himself into something he believed in. He had stood in front of the executive team in January and made a promise. The November 30th delivery was not just a date to him—it was a statement of what a disciplined, well-led team was capable of. Walking away from it, in his mind, was not a strategic adjustment. It was abandoning what they had started.

"That's what we committed to," he said, when she walked him through the findings. "We finish what we start."

She pushed back. He listened—genuinely listened—and returned to the same position with more force. "I've seen teams tear themselves

apart chasing the newest thing,” he said. “That’s how projects die—not from lack of ambition, but from lack of discipline. I’d rather deliver something real than chase a moving target in the final quarter.”

He was not dismissing the data out of carelessness. He was weighing it against what he saw as the greater risk: a team losing its focus too close to the finish line. His conviction was not for show—it was genuine. And that made it harder, not easier, to argue with.

She did not push back further. She returned to her desk. The rest of the committee received her debrief in silence.

What was happening inside Team Alpha’s leader was a trap: his identity had become so entangled with the original plan that abandoning it felt like personal failure. The new data was not the problem. The new data was, in fact, a gift. But the leader could not see it as a gift—he saw it as a threat. A threat to the discipline and focus he had worked so hard to build. In his mind, his primary duty to the team was to project conviction and steadiness under pressure. Entertaining the new data meant opening a door he had decided, on behalf of everyone, should stay closed.

“The hardest data to act on is the data that contradicts a decision you are proud of. That is also, almost always, the most important data you will ever receive.”

Nobody else was told about the technological shift and the insights gained by the committee. The leader absorbed the information and quietly set it aside. The original plan remained unchanged. The deadline stayed fixed. Olympus would ship on November 30th. The team continued building. But the next occasion on which the September

findings might have changed anything was already approaching: the Monthly Business Review, scheduled for the beginning of October.

The October MBR: A Missed Opportunity

Every team in the company runs a Monthly Business Review (MBR)—a recurring meeting where progress is reported upward and the executive layer gains the visibility it needs to make decisions. In healthy organizations, the MBR is more than a reporting ritual: it is the moment when reality meets the plan—where gaps are named, decisions beyond the team’s authority are surfaced, and leadership can actually influence.

The October MBR was, in theory, the mechanism that could have redirected Olympus in mid-flight. The September findings belonged in that room. The question of whether to continue was exactly what an MBR exists to resolve. The executive layer had the authority to approve a change of direction. A difficult conversation in October might have saved months of misdirected effort.

In Team Alpha’s world, however, the MBR had become something else entirely.

The leader’s genuine pride in his team—his desire to present their work with the rigour and polish it deserved, to demonstrate that this was a team that took its commitments seriously—had gradually transformed the monthly review into something staged: scripted, rehearsed, managed for impression. In his own mind, this was simply preparation and a commitment to excellence. But that staging gradually hardened into process, and the process developed its own gravity.

Weeks before the October MBR, the document cycle began. The first review—the pre-meeting for the pre-meeting for the actual meet-

ing¹—brought together the document owner and roughly ten team leads. They reviewed structure and framing.

The second review was broader. Anyone with a stake was invited. The document they received was usually half-formed—by design. The owner had learned through prior cycles that this round would generate conflicting feedback, so putting real effort in before that arrived rarely paid off. The room dutifully reviewed something that was not yet ready to be reviewed.

A third session followed, attempting to converge. Usually it did not. An emergency review was scheduled. Then, days before the MBR, came the final preparation review—the first session that included Team Alpha’s leader. He arrived with fundamental feedback. The document that had survived four rounds of review was substantially rewritten in the final hours. Further emergency sessions followed.

“When a team spends more energy on how the work looks than on how the work is going, it has crossed the line from results into optics. The crossing is always quiet.”

What made this remarkable was not the dysfunction—it was that no one questioned it. The reviews appeared on a fixed agenda distributed at the start of each month. The process had a name. It had calendar entries. It had become, in the fullest sense, official. And once a process is official, attending it becomes the work. The document owner was no longer trying to write a great document. The goal had shifted to simply completing the review cycle.

¹In a September 2024 memo, Amazon CEO Andy Jassy named this pattern explicitly, calling out “pre-meetings for the pre-meetings for the decision meetings” as a bureaucratic artifact that slows organizations down.

The costs were invisible until you added them up. Dozens of people, two to three hours each, across five or six sessions every month. This was not meeting overhead—it was the systematic extraction of the organization's most valuable resource: the focused attention of its most capable people, redirected from actual work toward the management of appearances.

The same instinct had spread beyond the documents. Engineers did not circulate half-formed ideas or rough implementations—there was no safe channel for it. Work, like the MBR document, was expected to arrive polished and defensible. Anything less invited the kind of scrutiny the team had learned to avoid.

The content of the documents themselves told a secondary story. In a healthy MBR, the document surfaces problems that need executive attention, risks that require decisions, and requests that only the leadership layer can grant. Team Alpha's MBR documents left these out. They were studies in carefully constructed reassurance: green metrics, calibrated language, and a narrative that positioned every challenge as already managed. The Olympus project, as described in those documents, was on track. The team was delivering. The deadline would be met.

None of this was technically false. It was simply, and consequentially, incomplete. The September findings did not appear. The pivot question was never surfaced. The executive layer, reading documents polished to flawlessness, had no reason to ask the questions that might have changed the outcome. They were, in the most literal sense, looking at a Potemkin facade—and seeing exactly what they had been shown.

The MBR that could have been the turning point of the Olympus project had instead reinforced its original direction. The leader left the room with a green scorecard and the quiet confidence that, for another month, the narrative had held.

The Team That Forgot to Look Up

The pattern had been visible throughout the year: no mechanism for revisiting whether earlier decisions were still valid, a leader who absorbed inconvenient data and set it aside, and a team whose calendar was already packed with meetings and contained no standing forum for questioning direction. The October MBR did not interrupt this trajectory—it quietly confirmed it.

From that point on, the focus narrowed further. Every conversation drifted toward execution. Every meeting became a check on progress, not purpose. And slowly, almost imperceptibly, the shared objective simplified into something unspoken but universally understood: deliver a green status by the end of November.

With no structured space for reflection or challenge, the “thinking layer” of the system had effectively disappeared. Questions about whether the work was still the right work faded first from meetings, then from side conversations, and eventually from individual consideration. What remained was execution.

By November, the team was moving quickly and quietly, delivering constantly, and keeping status green. But the underlying question—*are we still doing the right things?*—had stopped being asked altogether. Team Alpha had become pure delivery fixated: disciplined, focused, and efficient—locked onto execution as an end in itself. They were moving fast without noticing that they had stopped asking where they were going.

“The risk in high-pressure situations is rarely a lack of speed. It’s that when heads are down during execution, looking up becomes the first thing the team forgets.”

The One-Person Bottleneck

The missed opportunity of the October MBR had a second consequence—subtler, but just as damaging. By steering the discussion away from uncomfortable questions, the leader hadn't quieted his anxiety. He had deepened it. He had held the line once, but only because he'd been there to hold it. Beneath his outward conviction lay a quiet distrust—not of individuals, but of what might happen if he loosened control. So he didn't. That instinct infected everything from planning to deployment.

What followed felt, from the inside, like dedication. The leader inserted himself into everything—no architectural decision, integration approach, or cross-team commitment moved without his review. He worked long hours, tracked every thread, and stayed convinced his presence alone would carry the team through. Yet people needing decisions often found themselves waiting—sometimes hours, sometimes days—for his response.

The leader saw this not just as leadership, but as the strength and responsibility his role demanded. The team experienced it as a wall.

The most damaging effect wasn't any single delay—it was how the bottleneck reshaped the team's behavior. People stopped making decisions they once owned. The risk of acting without approval felt too high, so they waited. Over time, that waiting quietly eroded the team's sense of ownership.

The people closest to the implementation held the deepest knowledge of the system—knowledge that could have shaped a substantially better customer experience. But getting their insights to a place where they could influence decisions meant navigating an opaque chain of approvals with no clear entry point and no guaranteed outcome. Most stopped trying after one or two attempts.

What they did instead was execute: write the code they were told to write, integrate the systems they were told to integrate, and leave strategic questions to the one person authorized to answer them. This did not feel like surrender. It felt like reading the room: *this is what is expected of me here*. And so they stopped expecting more of themselves. They stopped questioning, anticipating, owning. A team that had been hired for its judgment was losing the habit of exercising it—and in doing so, becoming less capable than the sum of its parts.

Eyes on the Metric, Not the Customer

By late October, weeks of engineering effort had already been invested in an architecture the September findings had called into question. With every decision now flowing through a single point of approval, the team had lost the capacity to question where it was headed. What remained was the original plan, the original deadline, and the original metric.

In the planning documents, Olympus's success had been defined by a primary metric: recommendation accuracy of 30%—the share of cases in which the system's top recommendation matched what the user actually chose to engage with. The number had been chosen carefully in January, based on the data available then. It was a reasonable metric.

The science team had pushed hard: accuracy had climbed steadily through the autumn and now sat at 29.2%. But the gains had stalled. The latest rounds of tuning were no longer moving the metric in any meaningful way. Ideas for closing the gap existed—a fundamental change to the model's embedding architecture, an alternative approach to negative sampling—but each carried a significant cost: weeks of work and most of the team's remaining capacity.

The leader authorised the work. The reasoning was straightforward: the goal was 30%, and it had to be green. The work began immediately.

What was never asked was whether closing that final 0.8% would produce any meaningful difference for the customer. The team was heads-down, executing. The people who might have raised the question were occupied elsewhere—including preparing the next MBR document.

The final weeks of Olympus were consumed by the pursuit of a number. The customer disappeared from the conversation. Whether the thing being built was useful, whether it served the need it had been designed to serve, whether the energy being spent on accuracy optimisation might have been better directed at anything else—none of these questions had a place to be asked. Goodhart's Law² had done its quiet work: the metric had become the mission. The number had replaced the customer. The Olympus goal, on November 30th, would be green.

A team that loses focus on the customer optimises for the wrong outcomes. The moment you stop asking what value your metric represents, it starts driving you."

Nobody's Job

Four weeks before launch, a request arrived from the executive layer: a new companion feature—distinct from the core Olympus engine but closely related—needed to be ready alongside the main delivery. It was not a small ask—it would consume most of the remaining engineering weeks—but it was framed as important to the product narrative and to the story leadership planned to tell at the November delivery.

²Goodhart's Law, named after economist Charles Goodhart, holds that "when a measure becomes a target, it ceases to be a good measure." Originally an observation about monetary policy, it applies with equal force to any situation where the proxy metric for a goal becomes the goal itself.

The leader of Team Alpha responded with an email to the entire engineering team. The feature needed to happen, he wrote. It was a priority. Anyone who needed to deprioritise something else to make room for it was authorised to do so.

The team responded with diligence. A thread formed. People discussed resourcing. The feature was added to the sprint board. Capacity was notionally identified. All the motions of organised work were performed. And then everyone waited—quietly, reasonably—for someone else to pick it up. No one wanted to abandon what they were already deep into, so no one did.

What followed is a pattern so common it has a name: the Bystander Effect.³ No one had been told, explicitly, that this was their problem to solve—and their failure to own it if it went unsolved. Each engineer looked at the task, assumed someone better positioned would step forward, and returned to the work they already knew was theirs.

By mid-November, the feature had barely moved. When the leader asked for a status update, the answer from the thread was technically accurate: work was in progress, resourcing had been discussed, no blockers had been formally raised. What it did not communicate was that no one had truly started—because in Team Alpha, the status update had long since been decoupled from the underlying reality.

The feature did not ship. It was quietly removed from the November 30th scope in the final week, attributed in the internal documentation to “technical dependencies.” No one in the executive briefing was told the real story.

³The *Bystander Effect* describes the phenomenon whereby individuals are less likely to take action when others are present. In organisational contexts, it manifests as diffusion of responsibility: the assumption that someone else is handling it becomes, collectively, the reason no one does.

The Outcome

Team Alpha delivered Olympus on November 30th. The launch metric was green: recommendation accuracy had cleared 30%. The launch announcement had been rehearsed in three pre-reviews. The executive layer heard a clean, well-structured narrative of on-time delivery.

Within six weeks, the picture clarified. Customer engagement with the new recommendations was lower than the model had predicted—substantially lower. The accuracy metric was green; the customer behaviour was not. The feature was technically correct and experientially irrelevant. It answered the question the team had been asked in January, in an environment that had since moved on. In the intervening months, customers had developed expectations shaped by a shifting competitive landscape—the same September data Team Alpha had quietly set aside. No one said, *We knew this would happen*. They didn't have to.

The talented, driven people on Team Alpha were exhausted. They had given everything all year: late nights, weekends, their best ideas, their judgment. Not from the work itself. Exhaustion from work well done is a recoverable condition. They were exhausted from the staging: from the pre-reviews and the polished documents and the careful management of what was visible and what was not. From building something they had privately suspected since September was not the right thing. From the specific, numbing tiredness of doing good work in service of a goal they no longer believed in.

The departure of the first two high-performing engineers was noted in the quarterly talent review. The reasons given in exit interviews were careful and professional. The underlying dynamic, understood by everyone who had been in those rooms, was not spoken aloud.

TEAM BETA: THE INVENTION MACHINE

The Kitchen Conversation

Team Beta's leader didn't learn about the new findings in a formal briefing. She heard it in the kitchen, when a data scientist, coffee in hand, caught her between meetings. He had not prepared anything. "I don't know if this is relevant," he said, "but I've been reading about some new approaches in the recommendation space and I keep wondering whether any of it changes the picture for Olympus. Probably nothing. Just something that's been sitting in the back of my head."

He had no narrative and no suggested next steps. Just a half-formed thought and enough trust in his leader to say it out loud, unpolished. That kind of moment—when someone feels safe enough to share an unfinished idea without rehearsal or agenda—is the visible surface of a culture where people do not self-censor.

She asked a few questions, went back to her desk, and spent the next hour reading. Then she pulled science, engineering, and product leads into a room.

Her first words were not a verdict. They were a question.

"If this is right—and I think it probably is—what does it mean for what we are building in Olympus?"

The conversation that followed was anything but comfortable. It led them to a conclusion no one wanted to face: in light of the new data, the algorithm and architecture they had spent months building were the wrong approach. Not catastrophically wrong—it would still function, still ship, still likely meet the original metric. But with this new understanding, choosing not to act would mean knowingly delivering something inferior to their customers. The gap between what they

had committed to and what they now knew to be better was suddenly impossible to ignore. The only question left was what to do about it.

The leader of Team Beta had a principle she had articulated to her team long before project Olympus began: new knowledge is never a threat. It is the most valuable thing the environment can give you. A team that treats incoming data as friendly—that leans toward it rather than away from it—will always outperform a team that treats new information as a challenge to its existing commitments.

She called this, simply, “friendly data.” When new information arrived that contradicted the plan, the question was not “how do we defend the plan?” It was “what does this tell us?”

“In fast-changing environments, rapidly adapting to new information always beats rigidly executing plans made when you knew less.”

The team spent two days analyzing the implications in depth. By the end, the conclusion was as clear as it was uncomfortable: the right path was a full pivot to the new approach—not an adaptation of what they had built, but a near-complete restart of the algorithm and architecture. Much of their prior work would be rendered obsolete. The timing was equally clear: the original commitment to deliver by the end of November could no longer hold. The pivot would push the timeline back by at least two months.

Team Beta’s leader chose a path that required far more courage than staying the course. Rather than concealing the issue, she planned to transparently share the team’s thinking at the upcoming Monthly Business Review. She would make the case to delay the launch by two

months. Whether the organisation would reward that courage was another question entirely.

The Courage to Change Course

Walking into the MBR meant publicly acknowledging that the original commitment was being revisited. The plan the team had spent months executing was now the wrong plan, and the ask was not a minor adjustment but a restart. This proposal required vulnerability and invited an obvious question: if the original approach was wrong, why hadn't the team seen it sooner?

The MBR document laid out three things clearly: a sharp analysis of the new data, a specific description of what the better approach would deliver and why, and an honest accounting of what the change would cost. It did not minimize the cost. It did not pretend the transition would be seamless. Instead, it presented the tradeoff directly: the team would likely miss the November 30th deadline—the new estimate was end of January—but the product delivered then would create far more customer value than anything they could ship in November.

Nobody in the room was happy about the pivot. What followed was an intense, uncomfortable discussion. The entire 90 minutes of the Monthly Business Review were consumed by this single decision. Every assumption was challenged. The costs, risks, and reasoning were scrutinized from every angle. At one point, a sponsor asked, "So you're telling us we wasted the last six months?"

But the team came prepared: a data-driven argument, a clear customer-value narrative, and evidence they had done the hard thinking rather than the easy thing. They had shown transparency and vulnerability both in the document and in the room. So the debate, though heated, remained productive.

By the end of the session, the pivot was approved.

Clearing the Path

Before a single line of new architecture was written, Team Beta's leader gathered the team for a different kind of meeting. She knew that after any ambitious pivot, the next bottleneck would be decisions—the dozens of small and medium choices that needed to happen every day for execution to move at speed. In her experience, the most common source of delay in teams building complex technical products was not engineering capacity or unclear requirements. It was ambiguity about who was allowed to decide what.

Her response was to resolve that ambiguity before the work began. For every major category of decision the team would face, she wanted an answer to one question: at what level of authority does this decision sit?

The mapping was built in a two-hour session with the whole team. Every category was discussed until the level felt right—not comfortable for the leader, but right for the speed and quality of decision-making the project required. For instance, redefining the metrics used for evaluation required a group decision. But launching a small A/B test belonged to every individual on the team—no review, no approval, just the results to show for.

“Clarity about who decides is not bureaucracy. It is what turns capable individuals into a fast, effective team.”

By the end of the session, everyone knew which decisions they owned completely, which required a quick check-in, and which warranted a broader conversation before proceeding.

In the days that followed, decision after decision was made without confusion, without escalation, without anyone stopping to ask who was supposed to call it. Each person knew their authority and used it. The team moved with a clarity and confidence that had not been there before. Nothing was blocking them from using the judgment they already possessed.

Rough, Real, and Running

The world Team Beta was navigating did not reward exhaustive planning. In conditions of genuine uncertainty—shifting data, untested assumptions, a direction that had already been revised once—a plan that tries to anticipate everything before building anything is not rigor. It is delay dressed as diligence.

Team Beta's leader answered this directly: do not wait for the plan to be perfect. Build the smallest meaningful version as quickly as possible, and let reality start teaching you what the plan never could. Not a sketch on a whiteboard—a working prototype, running against real data, producing outputs you can actually evaluate. A minimum loveable product: not feature-complete, not production-ready, but real enough to tell you whether the direction is right. Their rhythm was simple: build something real by the end of each week, put it in front of test users, gather feedback, and decide what to change next.

“The fastest way to discover you are wrong is to build something real. That is also the fastest way to get right.”

Within three weeks of the pivot, Team Beta had a working prototype of the new architecture running against live customer data. It was not

clean. It was not documented. By any standard of software craftsmanship, it would need significant rework before it could be considered production-ready. But it answered the questions that mattered: was the new approach producing better recommendations? Was the customer experience what the data had predicted? Were there surprises in the implementation that the plan had not anticipated?

There were surprises. There always are. A critical API integration that had been assumed straightforward turned out to have a compatibility problem and required a different approach. The lead engineer who discovered it raised it in the team's daily standup without drama: "I don't have a solution for the API gap yet. I want everyone to know it's there. Does anyone have context on this layer?"

Two engineers from different parts of the codebase had relevant experience. The problem was solved in a day—not because the lead engineer was a uniquely creative problem-solver, but because she had said, clearly and without embarrassment, that she did not know the answer.

Every Week, Something Changed

The prototype was never meant to be the answer. It was meant to start the conversation.

Three weeks after the pivot, the first working version was shared with a small group of internal users across the company—other product managers, domain experts, and select leaders who could act as stand-ins for real customers. It was not a launch. It was a deliberately incomplete system placed in the hands of people who could stress-test assumptions and surface blind spots. The question was not whether the product was ready, but whether the direction was right.

Feedback came back within days. Some of it confirmed what the team had built. More of it did not. One pattern appeared across almost every response: users kept reaching for something the team had deliberately moved to a later milestone—a filtering capability they had deprioritised on the grounds that it was a secondary use case. The assumption had seemed reasonable in planning. In practice, it was wrong. Without it, the recommendation experience was technically sound and practically limited. Users could see what the product was trying to do; they just couldn't do what they actually needed.

The product lead flagged it at the next morning's standup. There was no drama in how she raised it—just the feedback, the pattern, and a question: do we stay on the current track, or do we pull this forward? The discussion was uncomfortable. Pulling it forward meant partially redoing work already completed. But the team had committed, when the pivot began, that feedback would drive the plan—not the other way around. Within the hour, the priority had shifted.

The next iteration shipped the following week. The feedback that came back confirmed the call.

None of this was treated as a setback. The team absorbed each round of feedback, triaged what mattered most, and shipped an updated version before the next review cycle began. Then they did it again. New version out, feedback in, changes made, shipped. As many times as the timeline allowed, at the fastest pace the team could sustain.

From the outside, the cadence looked almost relentless. Inside the team it felt less like pressure and more like momentum. Each cycle left them with fewer unknowns. Features that had seemed uncertain were clarified by feedback that no internal review would have caught. The team did not slow down to process what they were learning—they had built the processing into the rhythm of the work itself.

A Mid-Level Engineer Takes Ownership

When the companion feature request arrived in mid-October—the same request that would lead Team Alpha in the Bystander Effect—Team Beta’s leader did not start where most leaders would.

She did not open a spreadsheet. She did not scan the team for spare capacity or calculate who still had headroom. Instead, she ignored re-sourcing entirely and asked a different question: who on this team, given complete focus and full accountability, could be stretched into something they had not done before?

The answer was a mid-level engineer—talented, currently working at about seventy percent of his potential on work he had largely mastered. The companion feature was adjacent to his existing skills but meaningfully more ambitious. The leader spoke to him directly: “I want you to own this. Not as one of your responsibilities. As your only responsibility for the next six weeks. You have full authority over the architecture, the timeline, and the delivery approach. The decisions are yours.” She paused. “This is a stretch. I think you’re ready for it. Tell me what you need—including people. Whoever you ask for will be pulled from their current work and assigned to this one hundred percent. No split focus—yours or theirs.”

“The question is never who has capacity. It is who is ready to grow. Those are rarely the same person.”

The engineer’s response, in the retelling, was a grin and a single line: “So you’re saying I should cancel all my other meetings?”

He meant it as a joke. But it captured something real: the weight of a challenge that demanded total engagement. He was not overwhelmed.

He was ignited. He was outside his comfort zone every single day. But he had never been more engaged. Ownership turned the stretch into a challenge he wanted to win.

The companion feature shipped on time, and it shipped well. In the January presentation to the executive team, it received the strongest response of anything on the agenda.

The Outcome

Team Beta missed the November 30th deadline. The conversation with the executive layer about the January delivery had been honest, data-driven, and difficult. But when the product launched in January, customers responded with something teams chase for years: not incremental improvement, but a disruptive shift. The excitement was real. A product manager forwarded an email from a user who wrote, “I actually look forward to see these recommendations now. What changed?”

The executive sponsor, who had approved the October pivot with reluctant confidence, presented the January results to the broader company leadership. The narrative was not “we missed November.” The narrative was “we made the right call when the data told us to.” Team Beta’s leader had understood, from the beginning, that the reputation worth protecting was not the reputation for delivering on a specific date. It was the reputation for building things that worked.

The team began the new year with something Team Alpha had lost and would struggle to recover: trust. Trust from the executives who had seen them make a hard call and back it with data. Trust among themselves—the specific, durable variety that comes from having built something genuinely good together, under real pressure, with the freedom to make real decisions. And trust, most importantly, from their customers.

THE MECHANISMS THAT MADE THE DIFFERENCE

The divergence between Team Alpha and Team Beta was not a story about talent, ambition, or resources. Both teams wanted Olympus to succeed. Both leaders deeply cared about their people and their product. The difference was in the culture—in the specific, repeatable practices that each team had built, or failed to build, long before the September fork arrived.

Here are the concrete practices—the mechanisms—that produced Team Beta’s outcome. Any leader or team can build them.

1. Protect Time for Brainwork

The most quietly destructive decision a team can make is to fill every hour with execution—Grindwork. Grindwork is the execution of agreed-upon tasks: the phase of a project where the plan is set and the work is delivered. It includes the day-to-day flow of tasks—measurable, visible, and satisfying to report. It’s crucial for driving progress, yet it can also lead to tunnel vision.

By contrast, Brainwork is stepping back to reflect, question assumptions, and ask whether you are still pursuing the right priorities. Brainwork is time to think—truly think, without producing a deliverable. It produces nothing that appears on a sprint board. In any week with real delivery pressure, it feels like a luxury.

But Brainwork is never a luxury. It is the mechanism that keeps execution pointed in the right direction. It is the only way to avoid the Day 2 trap,⁴ where teams become so immersed in operations that they

⁴The *Day 1 / Day 2* concept comes from Jeff Bezos’s 2016 shareholder letter. Day 1 reflects a startup mindset—fast, customer-focused, and experimental—while Day 2 marks the drift into process, bureaucracy, and eventual decline.

lose sight of their purpose. Without Brainwork, you get stuck doing the same things without remembering why. You continue down a losing path simply because you are too busy executing to pause.

Why is Brainwork so hard to protect? Grindwork is easy—it requires less brainpower, it is visible, and it is self-reinforcing. Completing action items shows progress. Blocking an hour to think shows nothing. And as humans, we have a powerful tendency to do the easy thing—especially when we are stressed. Think of the last time an important, complex task was waiting for you, yet you spent hours answering unimportant emails instead. The faster you execute, the more you feel you cannot afford to stop—especially in a fast-paced, ambiguous world where it is easy to forget what you set out to achieve.

The solution is a standing calendar commitment: time blocked at a regular cadence for thinking, not doing. Not a status meeting. A dedicated forum where the team asks uncomfortable questions: Are we still building the right thing? What has the environment told us? What do we know now that we didn't know when we made this plan?

This time will feel wasted in months when the plan is correct. It will feel invaluable when it is not. The September pivot was not a crisis for Team Beta because they had already built the habit of asking the hard questions. By the time the new data arrived, asking those questions happened naturally. That is how you stop grinding past decisions without purpose and start building what matters.

2. Map Decision Authority

Ambiguity about who decides is one of the most expensive and least-discussed sources of organizational drag. When decision authority is not explicit, work slows as people hesitate, defer, or assume someone else will act—an organizational bystander effect. Meanwhile, capable people

either hold back from using their judgment or overstep their mandate. Both paths create the same outcome: frustration, friction with leaders, and tension over who was actually supposed to decide.

The solution is simple: before work begins, map every major type of decision to an agreed authority level. Think of a spectrum with two extremes. At one end, the leader decides alone and simply informs the team. At the other end, the team decides fully, sometimes without even reporting back. Between these extremes are intermediate stages where decisions are progressively shared through consultation, joint ownership, and delegated authority with oversight.⁵

The real power is not the map itself. It is the transparent conversation required to create it. That conversation—about who owns what, where the leader’s judgment adds genuine value, and where it becomes a bottleneck—is the most productive authority discussion a team can have. It surfaces untested assumptions and clarifies expectations that the leader and the team have been holding separately—and usually differently.

Map it collaboratively. Revisit it when the work changes. Make it visible so that for any given decision, everyone knows where they stand.

3. Build, Learn, Iterate

When the future is hard to predict, the right response is not a longer planning phase, but a faster learning loop. Product development should operate as a continuous flywheel: build, share, learn, iterate. Speed of learning matters more than initial correctness.

This reframes the question of planning horizons. Traditional ap-

⁵Jurgen Appelo defines seven levels of delegation, from “Tell” (leader decides alone) to “Delegate” (team decides fully), in “Management 3.0: Leading Agile Developers, Developing Agile Leaders.”

proaches often emphasize long-range plans—three-year strategies—but these rarely survive contact with reality. A more practical guideline: be precise about the next month, directional about the next three, and deliberately vague beyond that. In volatile environments, the value of a plan decays faster than the effort to produce it.

Start by applying the Pareto principle:⁶ Identify the small fraction of functionality most likely to generate meaningful value. Build it quickly in the real environment where it will run. The goal is not completeness, but contact with reality as early as possible.

Then immediately share it. This takes courage. Showing something unfinished feels risky. But the real risk is waiting until it is too late to learn you were wrong. Cunningham's Law reinforces this dynamic: the fastest way to improve something is not to refine it in isolation, but to expose an imperfect version and invite correction.⁷ A polished product tends to suppress feedback; a rough prototype makes it easier for others to react, challenge assumptions, and surface issues you did not anticipate.

From there, the key is rapid iteration. After each feedback cycle, ask three questions: what will we change? What will we not change? And what do we now know that we didn't know before? Without this closure step, feedback accumulates without convergence; teams collect insights endlessly but fail to turn them into directional learning. When this step is explicit, each loop tightens the model of reality.

Over time, this creates a compounding cycle: build → share →

⁶The *Pareto principle* (80/20 rule) holds that a small fraction of causes—roughly 20%—produces a large majority of effects—roughly 80%. Not all inputs are equal; a few things matter much more than the rest.

⁷*Cunningham's Law* is attributed to Ward Cunningham, creator of the first wiki: "the best way to get the right answer on the internet is not to ask a question but to post the wrong answer."

learn → decide → adjust → rebuild. The system improves not through single large leaps, but through many fast, feedback-driven cycles that progressively eliminate uncertainty.

4. Single-Threaded Ownership

Every piece of work that matters should have one name on it.

Not “the team is responsible” or “this falls under the platform pillar.” One person. One primary accountability for a defined period. And the authority to make the decisions required to deliver it. This is single-threaded ownership: the deliberate assignment of end-to-end responsibility to prevent the diffusion of accountability that leads to delay, indecision, and the bystander effect. This feels risky to leaders trained to spread risk. But shared ownership is not insurance against failure—it is a guarantee of confusion.

When ownership is single-threaded, several things change. Decisions happen faster, because one accountable person can make them without defaulting to coordination. Problems surface earlier, because there is no ambiguity about who owns them. And the quality of the work improves, because complete ownership—held by one person rather than shared across many—creates a level of focus and care that collective responsibility rarely produces.

Single-threaded ownership also sharpens incentives. When responsibility is diffuse, success and failure are equally easy to misattribute. When it is clear, outcomes become legible: progress, delays, and trade-offs all have a visible owner.

The owner should be chosen with growth in mind. The strongest assignments stretch the individual: ambitious enough to require full engagement, bounded enough to remain achievable, and supported

enough that the stretch does not become a trap. Given the choice between assigning critical work to an overloaded expert or to an energized person ready for the next step, choose the latter. Focus, ownership, and motivation beat experience alone.

5. Make Every Meeting Earn Its Purpose

The most common failure mode in meetings and reviews is not inefficiency—it is the absence of value. Many meetings simply do not change anything: they consume time, create the appearance of progress, and end with no decisions, no surfaced risks, and no meaningful outcomes. It is no surprise that meetings have earned their reputation over time as a tax on real work.

That conclusion is wrong. When designed well, meetings are one of the most powerful mechanisms an organisation has—where decisions are made, alignment is created, and trade-offs are resolved. Their poor reputation has a simpler cause: most meetings are set up and run incorrectly.

Four mechanisms help keep meetings and reviews pointed at what matters.

Use a narrative as the guiding artefact. Begin with a concise written narrative that makes the objective explicit: what are we trying to achieve, what are the key questions, and where is input required. A well-crafted narrative aligns the room before discussion begins, ensures shared context, and forces clarity upfront. It keeps the conversation anchored on outcomes rather than drifting across disconnected points.⁸

State the intent explicitly, and make it specific. A monthly

⁸Amazon is known for its narrative-driven meeting culture, where concise written documents replace slide decks and serve as the basis for discussion. The document itself is not the goal; it is a mechanism to clarify thinking and enable better decisions.

business review, for example, should not be a tour of metrics. Its purpose is to surface what the team is most stuck on: decisions that need leadership input, risks that must be named in front of executives, and hotly debated topics that require resolution. State this clearly in the shared agenda and the narrative. Then call it out again at the start of the meeting. This changes what people focus on—and what the room considers a successful outcome.

Invite only those with a direct contribution to make. Every extra person who does not need to be there shifts the dynamic away from candid discussion. At scale, frank debate becomes unlikely—few people will openly admit uncertainty or gaps in front of a large audience. Worse, in a big room, individuals optimize for appearance over honesty, trying to impress rather than solve. For each person on the invite list, ask one question: will they contribute something the room cannot produce without them? Or are they there because leaving them out would feel awkward? The second answer is not a reason.

Close every meeting with explicit next steps, owners, and an honest commitment to follow up. Discussed items that produce no next steps have, in practice, produced nothing. For each action agreed in the room, there should be a named owner and a date. Those actions should be shared back to all relevant stakeholders after the meeting ends—not buried in notes that no one reads, but surfaced explicitly in the channels where accountability lives. The test of a meeting's value is not what was said in the room. It is what changed because of it.

A note on moderation. Stronger moderation is often cited as the key to more effective meetings. It is, at best, a symptom-level fix. It treats drift as a facilitation problem, when the deeper cause is almost always unclear purpose—or the wrong people in the room. When a team has a clear intent and the right attendees, the meeting largely self-moderates.

Investing in a skilled moderator before fixing the invite list and the intent is solving the wrong problem.

6. *Counter Escalation of Commitment*

There is a force in every organization that resists change not because change is wrong, but because momentum already exists. Psychologists call this escalation of commitment: continuing to invest in a course of action even when evidence for it has weakened or disappeared entirely.⁹

Stopping or re-evaluating something that has been in motion for a long time is difficult. People struggle to accept what feels like wasted effort—the time and energy already invested become losses. That loss stings even more when it means letting go of something you helped build and may still be proud of. Then there is habit: what people have done for months, what they have learned to navigate, and what they now know deeply. Changing course means leaving your comfort zone and stepping into uncertainty.

At its core, this is inertia: the tendency to keep moving in the same direction simply because you are already moving that way. Without a conscious force to change course, momentum carries you forward—even when the path no longer fits.

You have almost certainly heard the tell-tale phrase: “we have always done it this way.” It appears in conversations about processes, tools, meeting structures, project approaches, and product decisions. It is almost never a reasoned argument. It is the sound of escalation of commitment speaking—the accumulated weight of past decisions presenting itself as a justification for future ones.

⁹*Escalation of commitment* was formalized by Barry Staw in his 1976 paper “Knee-deep in the Big Muddy,” describing the tendency to double down on a failing course of action to justify prior investments.

The antidote is not more willpower. It is a structural habit: schedule dedicated, recurring reviews whose only question is “should we still be doing this?” Not a general retrospective. A focused audit of specific investments—projects, features, workstreams, and, especially, meetings. Each item gets a simple test: if we were not already doing this, would we choose to start? If the honest answer is no, stop.

Apply this to your project portfolio. Apply it to the feature list. Apply it to your meetings—and with particular discipline, to recurring meetings. Recurring meetings have a unique tendency to persist long after their value has passed. Cancelling a meeting requires someone to act. Continuing it requires no one to do anything at all.

The organisations that perform best in volatile environments are not the ones that execute their plans most faithfully. They are the ones that can stop, examine what they are doing, and discard what no longer serves them—before the environment forces the decision on them.

THE CHOICE THAT DEFINES THE DELIVERY

The Olympus story ends in two very different places. Team Alpha delivered on time and shipped a product that missed its customers, exhausted its people, and left the organisation wondering, quietly, how a project that had looked so green in every review had produced so little value in reality. The talented engineers who had built it began, one by one, to direct their energy elsewhere. The leader who had held the line so tightly found himself holding nothing at all.

Team Beta missed a deadline, changed a plan, admitted a gap, shared scrappy prototypes, and delivered a product that their customers actually used—one built by a team that ended the year stronger, more trusting,

and more capable than when it had started.

Same goal. Different choices. Entirely different products.

The most important insight is not that Team Beta had better people or a wiser leader. The mechanisms were in place before the September data arrived. Defining authority levels in advance meant that decisions were being made at the right level before the pivot required fast ones. The habit of sharing early and asking “what are we missing?” meant that the September data was treated as useful rather than threatening. The discipline of running meetings with a clear purpose meant that the right conversations happened rather than the comfortable ones. And the practice of protecting time for Brainwork meant that when the environment moved, the team had the ability to notice.

Good intentions don’t build resilient teams. Mechanisms do.

We are operating in a world where the environment will always move faster than the plan. The AI era will not slow this down. The systems being built are more capable, more interconnected, and more sensitive to context than anything that came before. The gap between what we committed to in January and what the customer needs in September will remain wide, and will require leaders and teams willing to close it honestly rather than paper over it.

What we can control is not whether the environment moves. What we can control is whether our teams have built the habits—the real ones, the practiced ones—to move with it. That starts with the meeting you protect for thinking rather than reporting. With the authority levels you define before the pressure is on. With the scrappy document you share before you are ready. With the recurring question you schedule into the calendar and ask without flinching:

Build for the customer. Stop what no longer serves. Own it completely.

The delivery will follow.

CHAPTER SUMMARY: PROCESS VS. PURPOSE

	Team Alpha <i>The Potemkin Grind</i>	Team Beta <i>The Invention Machine</i>
Mindset	Escalation of Commitment: “We agreed to this.”	Counter inertia: ask “should we still be doing this?”
Work Style	Pure Grindwork: execute the plan blindly.	Integrate Brainwork: think, reflect, and pivot.
Decisions	Leader bottleneck: nothing moves without approval.	Decisions mapped explicitly in advance.
Reviews	Prep-to-the-Prep-to-the-Prep: polish the facade.	Clear intent, explicit next steps.
Development	Plan exhaustively, move slowly.	Rapid prototyping: build the vital 20% first.
Feedback	Protect work until it is polished.	Share early and scrappy.
Uncertainty	Showcase confidence, suppress the unknown.	Name the gap openly, invite answers.
Ownership	Bystander Effect: add it to the pile.	Single-Threaded Leadership: one name.
Success Metric	Make the metric green.	Make the product useful.
Outcome	Product missed customers; talent departed.	Flywheel of trust; team energised and growing.

SIX MECHANISMS FOR A PURPOSE-DRIVEN TEAM

1. Protect Time for Brainwork

Situation — A calendar full of delivery meetings leaves no space for strategic reflection—the team executes the plan without asking whether the plan is still right.

Behaviour

- Block a standing, recurring time slot dedicated to strategic reflection—separate from delivery rituals
- Ask: are we building the right thing? What has the environment told us since we last checked?
- Protect the slot when delivery pressure mounts—that is precisely when it matters most

Impact — Teams that never pause to look up stay productive until the moment they arrive at the wrong destination. Protected thinking time keeps execution pointed at what the customer actually needs.

2. Map Decision Authority

Situation — A significant decision needs to be made—and it is unclear who has the authority to make it.

Behaviour

- Before work begins, map every significant decision category to an agreed authority level—from leader decides alone to team decides independently

- Build the mapping collaboratively; revisit it when the work changes
- When a decision falls outside the map, clarify the level before acting, not after

Impact — Decision bottlenecks are not a capacity problem—they are a clarity problem. A shared authority map lets capable people move at speed without waiting for permission they should already have.

3. Build, Learn, Iterate

Situation — The team is building under uncertainty—requirements shift, assumptions are untested, and the right solution is not yet clear.

Behaviour

- Build the vital 20% first: fast, rough, in the real environment—treat it as a question, not a deliverable
- Share it before it feels ready; a rough prototype invites honest correction, a polished one invites polite approval
- After each feedback cycle, ask three questions: what will we change? What will we not change? What do we now know that we didn't before?
- Repeat—build → share → learn → iterate—as many times as the timeline allows

Impact — Speed of learning matters more than initial correctness. Each cycle eliminates unknowns and tightens your model of reality. The team that iterates fast doesn't just build better—it removes the uncertainty that slower teams are still managing at launch.

4. Single-Threaded Ownership

Situation — An important initiative has multiple stakeholders, contributors, or sponsors—and no single person who owns the outcome.

Behaviour

- Assign one person to own one important thing—completely, not as one priority among many
- Choose the owner to stretch and grow them, not merely to fill a slot
- Make the ownership explicit and public: one name, one outcome, no ambiguity

Impact — Shared ownership is no ownership. When the result belongs to everyone, it belongs to no one—and when the Bystander Effect takes hold, the work drifts. One named owner changes the dynamic entirely.

5. Make Every Meeting Earn Its Purpose

Situation — A meeting is scheduled—whether a status update, a review, or a recurring sync—and its purpose is assumed rather than stated.

Behaviour

- State the intent of every meeting explicitly before it begins—the goal is always to achieve something, never to produce a good document
- Invite only those with a direct contribution to make; every extra person shifts the dynamic from discussion to performance
- Close with named owners, dates, and next steps—shared back to

all relevant stakeholders

Impact — A meeting without a decision is a rehearsal. The moment output is expected, the room shifts from performance to work. What is discussed without a named next step is rarely acted upon.

6. Counter Escalation of Commitment

Situation — A project, feature, or process has been running long enough that stopping it feels like failure—regardless of whether it still earns its place.

Behaviour

- Schedule recurring “should we still be doing this?” reviews—for projects, features, and meetings
- Apply the test: if we were not already doing this, would we choose to start?
- Stop what no longer earns its place before the environment forces the decision

Impact — The sunk cost is not the danger—the belief that spending more will recover it is. Every month an obsolete initiative runs, it consumes the energy that a new one needed. Stopping on purpose is not failure; it is the precondition for what comes next.